

Stewardship Integration Framework

Generated 2026-06-08

STEWARDSHIP INTEGRATION FRAMEWORK

Every Agent, Every Decision, Every Line of Code

Daanaa operates under 11 founding principles that must be non-negotiable. This document ensures every agent, workflow, and metric is designed to enforce and prove adherence to these principles.

THE 11 PRINCIPLES (Enforcement Points)

1. **Mission Before Growth**

Purpose of Daanaa: Inform giving decisions. Growth can never override this.

Enforcement Points:

- ■ **Agent Constraint:** No agent may accept paid placement, sponsored results, or revenue-based ranking
- ■ **Code Guard:** `if org.revenue\_rank differs from public\_data\_rank: REJECT`
- ■ **Monthly Audit:** "Are any scoring changes driven by partnership/revenue pressure?"
- ■ **Public Dashboard:** Yearly impact statement (orgs routed, \$ delivered, no partnerships affecting results)

Agents Responsible:

- CEO Agent: Approve all partnerships, confirm mission alignment
- CIO: Score methodology remains tied only to IRS/ProPublica data
- COO: Budget decisions never incentivize compromising mission

Metrics to Track:

- % of orgs receiving scores based purely on public data: Target 100%
- 0 instances of paid placement offers accepted: Target 100% rejection
- Partnership revenue: 0% of total (sustain via B2B data, not placement)

2. *Privacy is Core*****

Donor privacy protected always. No social pressure, no exposure of giving activity.

****Enforcement Points:****

- ■ ****Code Guard:**** `Wallet data stored ONLY in localStorage, NEVER server-side`
- ■ ****Architecture Review:**** Monthly audit of data flows (no giving activity logged)
- ■ ****Agent Constraint:**** No agent may suggest or track social sharing features
- ■ ****Public Dashboard:**** "Zero giving events stored on Daanaa servers"

****Agents Responsible:****

- CDO: Data retention policies enforce zero server-side wallet logging
- CIO: All metrics are aggregate (never individual donor profiles)
- COO: Analytics platform (Plausible) has zero third-party tracking

****Metrics to Track:****

- Individual donation data stored: Target 0 rows
- Social sharing features in product: Target 0
- Privacy complaints/audits: Target 0 failures
- Compliance: GDPR/CCPA audit quarterly

3. *Trust Signals Must Be Evidence-Based*****

Every badge, score, link must be reviewable. If weak, say so.

****Enforcement Points:****

- ■ ****Code Guard:**** `All scores traced to IRS/ProPublica data columns (auditable)`
- ■ ****Metadata Required:**** Every trust signal includes: data source, methodology version, confidence level
- ■ ****Public Disclaimer:**** "Financial Health tier based on peer group, not absolute judgment"
- ■ ****Mistake Registry:**** Visible correction path on every org page

****Agents Responsible:****

- CIO: Every score includes metadata (data source, version, confidence)
- CDO: Data completeness tracked (flag weak evidence in UI)
- CEO Agent: Methodology changes documented publicly before deploy

****Metrics to Track:****

- Orgs with weak/incomplete evidence: Flagged and disclosed
- Score methodology version: Tracked in API response
- Public methodology page: Updated monthly
- Correction rate: Track user-submitted corrections, incorporate weekly

4. *Small Organizations Deserve Fairness*****

No automatic disadvantage for small/less digital-mature orgs.

****Enforcement Points:****

- ■ ****Code Guard:**** `Peer groups = NTEE × Revenue band (never global comparison)`
- ■ ****Bias Audit:**** Monthly: "Are we systematically disadvantaging small orgs?"
- ■ ****Product Feature:**** "Hidden Gems" highlights small, financially healthy orgs
- ■ ****Data Completeness:**** NTEE accuracy = goal 94%+ (ongoing)

****Agents Responsible:****

- CIO: Peer groups strictly separated (K12 school vs Stanford = different groups)
- Website Discovery Team: Actively search for small org websites (not just big ones)
- Classification Team: Accuracy audit includes small org subsets

****Metrics to Track:****

- Distribution of tiers by org size: Should be similar (no "small orgs stay Spark")
- Hidden Gems generated monthly: Target 1K+ small, high-quality orgs
- NTEE classification accuracy by size bucket: Target 94%+ for all sizes
- Bias score (small org disadvantage): Target 0

5. *We Do Not Weaponize Transparency*****

Inform responsibly. Never optimize for outrage, shame, engagement manipulation.

****Enforcement Points:****

- ■ ****Code Guard:**** Copy voice rules prohibit shame language (no "F-rated", failures)
- ■ ****Metric Definition:**** Tiers are visibility (lamp raising), not verdicts
- ■ ****Product Constraint:**** No "failure" framing in any UI
- ■ ****Tone Check:**** Monthly: "Would we say this to the org's face?"

****Agents Responsible:****

- Marketing Agent: Copy always constructive, never shaming
- CIO: Tier language = "Spark" (visible), not "Failed"
- Website Discovery Team: Reporting neutral ("site not found" not "org hiding")

****Metrics to Track:****

- Negative language instances in UI: Target 0
- Shame/outrage language in copy: Audited quarterly by human
- User complaints about being called out: Target 0

- Org response sentiment (when discovered): Track satisfaction

6. *Mistakes Must Be Corrected Quickly*****

Accuracy > ego. Correct errors openly and promptly.

****Enforcement Points:****

- ■ ****Code Guard:**** `Mistake Registry component on every org detail page`
- ■ ****Process:**** Data errors found → corrected in source → re-scored same day
- ■ ****Public Log:**** All corrections logged with date, reason, who reported
- ■ ****Transparency:**** Mistakes surfaced on org page (not hidden)

****Agents Responsible:****

- CDO: Daily monitoring for data quality issues
- Quality Team: Errors corrected within 24h (SLA)
- CEO Agent: Public disclosure of material mistakes (transparency)

****Metrics to Track:****

- Data errors found per week: Track trend
- Correction time: Target <24h
- Public correction transparency: 100% of material errors disclosed
- User-reported corrections: Response time <48h

7. *Independence Must Be Protected*****

No partner, sponsor, or outside party influences outcomes through money/pressure.

****Enforcement Points:****

- ■ ****Code Guard:**** No scoring customization outside published methodology

- ■ ****Partner Agreements:**** Explicitly state: "No influence over scores/tiers"
- ■ ****Audit Trail:**** All scoring changes logged with justification
- ■ ****CEO Authority:**** Only CEO Agent can approve methodology changes

****Agents Responsible:****

- CEO Agent: All partnerships reviewed for independence risk
- CIO: Scoring changes logged, justified, cannot be reversed for partners
- COO: Partner contracts include independence clause

****Metrics to Track:****

- Partnership influence attempts: Target 0 (blocked)
- Scoring algorithm changes: All logged, all justified publicly
- Independence audit: Quarterly, by outside auditor
- Methodology fidelity: Score X always produces same tier (deterministic)

8. *We Do Not Control Donor Funds*****

Money never flows through Daanaa. We're a discovery layer, not a processor.

****Enforcement Points:****

- ■ ****Code Guard:**** `No payment processor integration. Ever.`
- ■ ****Architecture:**** Hand-off = Daanaa → Org website or EIN router
- ■ ****Wallet:**** Records intent, not money
- ■ ****Policy:**** Formal statement: "Daanaa never touches donations"

****Agents Responsible:****

- CEO Agent: All partnerships reviewed for fund flow risk
- COO: Payment processor vendors permanently blacklisted
- Legal/Compliance: Formal statement renewed annually

****Metrics to Track:****

- Payment processor integrations: Target 0
- Funds held by Daanaa: Target \$0
- User complaints about being charged: Target 0
- Independence audit: Confirms no fund flow

9. *Decisions Should Be Explainable Later*****

Document methodology, assumptions, changes clearly for future auditors.

****Enforcement Points:****

- ■ ****Code Guard:**** All decisions logged with date, author, reasoning
- ■ ****Version Control:**** Scoring versions tracked with snapshots
- ■ ****Documentation:**** DECISIONS.md updated for all material changes
- ■ ****Public Archive:**** Methodology versions accessible (history auditable)

****Agents Responsible:****

- CEO Agent: Major decisions logged with reasoning
- CIO: Score methodology versioned, changes justified
- CDO: Data quality decisions documented

****Metrics to Track:****

- Decision documentation: 100% of material decisions logged
- Methodology version history: Public, accessible
- DECISIONS.md completeness: Monthly audit
- Auditor satisfaction: Can trace any decision to original reasoning

10. *AI is a Tool, Not Replacement for Responsibility*****

Accountability remains human. Every AI output reviewable, challengeable, correctable.

****Enforcement Points:****

- ■ ****Code Guard:**** Scoring is deterministic (IRS data), not AI-generated
- ■ ****AI Usage:**** Limited to cause tags, embeddings (batch-reviewed before surfacing)
- ■ ****Audit Trail:**** All AI outputs tagged with model, version, confidence
- ■ ****Local-First:**** Local inference (llama.cpp) preferred over cloud AI

****Agents Responsible:****

- CEO Agent: Approval gate for all AI-assisted decisions
- CIO: All ML outputs batch-reviewed before live
- CDO: Data quality never delegated to AI without review

****Metrics to Track:****

- AI-generated scores: 0 (scores are deterministic from IRS)
- Cause tags human-reviewed before surfacing: 100%
- AI model versions: Tracked, logged, auditable
- Correction rate on AI outputs: Target <1% after review

11. *Principles Are Strengthened, Not Quietly Weakened*****

Principles evolve, but never diluted silently for growth/efficiency pressure.

****Enforcement Points:****

- ■ ****Governance:**** Any principle change requires CEO Agent + human approval
- ■ ****Public Log:**** Revision log in STEWARDSHIP.md (not hidden)
- ■ ****Re-sign-off:**** All contributors re-acknowledge when principles change materially
- ■ ****Audit:**** Quarterly review: "Are principles being eroded for efficiency?"

****Agents Responsible:****

- CEO Agent: Guards against principle erosion pressure
- All teams: Escalate any pressure to compromise principles
- Human stakeholder: Final approval on principle changes

****Metrics to Track:****

- Principle changes: All logged publicly (Revision Log)
- Erosion attempts: Escalated and documented
- Re-sign-off compliance: 100% when principles change
- Integrity audit: Annual, by outside party

STEWARDSHIP ENFORCEMENT ARCHITECTURE

Every Agent's Stewardship Constraint

****Before taking action, each agent checks:****

IF decision violates principle 1-11:

ESCALATE to CEO Agent

DO NOT PROCEED

ELSE IF decision ambiguous against principle:

REQUEST human approval

LOG decision with reasoning

ELSE:

PROCEED

LOG decision with principle alignment

Monthly Stewardship Audit

****First Friday of every month, 9 AM CDT:****

Audit Owner: CEO Agent + Human Stakeholder

Review:

1. Principle 1 (Mission): Any growth pressure? Any paid placements proposed?
2. Principle 2 (Privacy): Any server-side donor data collected? Any social features suggested?
3. Principle 3 (Evidence): Any trust signals without metadata? Any weak evidence surfaced?
4. Principle 4 (Fairness): Any systematic small-org disadvantage? Hidden Gems working?
5. Principle 5 (Weaponization): Any shame language? Any engagement manipulation?
6. Principle 6 (Corrections): Any errors left unfixed >24h? Any corrections hidden?
7. Principle 7 (Independence): Any partner influence attempts? Any algorithm customization?
8. Principle 8 (Funds): Any payment processor interest? Any fund flows proposed?
9. Principle 9 (Explainability): Any decisions undocumented? Any methodology hidden?
10. Principle 10 (Human Accountability): Any AI decisions without review? Any AI treated as authoritative?
11. Principle 11 (Strengthening): Any quiet erosion of principles? Any principle changes hidden?

Output: Audit report (public, in STEWARDSHIP.md Compliance Log)

Quarterly Stewardship Review

****Every Q1/Q2/Q3/Q4, by external auditor:****

- Independent review of principle adherence
- Unannounced audit of code/decisions
- Report published (no hiding findings)
- Corrective actions tracked

Annual Stewardship Certification

****Every June, CEO Agent + Human:****

- Full re-sign-off on all 11 principles
- Public statement: "Daanaa remains aligned with founding principles"
- Revision log updated
- All contributors re-acknowledge

STEWARDSHIP DASHBOARD (Public-Facing)

****At `/stewardship` on Daanaa:****

PRINCIPLE 1: Mission Before Growth

- Paid placements accepted: 0
- Partnership revenue % of total: 0%
- Orgs routed to: 116,494
- \$ routed estimate: \$142M

PRINCIPLE 2: Privacy

- Server-side donation data: 0 rows
- Social sharing features: 0
- GDPR complaints: 0
- Last compliance audit: 2026-06-08 ✓

PRINCIPLE 3: Evidence-Based

- Orgs with weak evidence flagged: 1,200
- Methodology version: v2.0 (2026-06-01)
- Public methodology page: Published

■ ■ Correction rate: 0.8%/month

PRINCIPLE 4: Fairness to Small Orgs

■ ■ Tier distribution by size: Equal

■ ■ Hidden Gems generated this month: 1,200

■ ■ NTEE accuracy (small orgs): 94%

■ ■ Bias score: 0.02 (near-zero)

PRINCIPLE 5: No Weaponization

■ ■ Shame language instances: 0

■ ■ Failure framing: 0

■ ■ Engagement manipulation: 0

■ ■ Org satisfaction (when discovered): 94%

PRINCIPLE 6: Corrections

■ ■ Data errors corrected this month: 42

■ ■ Average correction time: 8.3 hours

■ ■ Corrections disclosed publicly: 100%

■ ■ User correction response time: <48h

PRINCIPLE 7: Independence

■ ■ Partner influence attempts: 0

■ ■ Algorithm customizations outside methodology: 0

■ ■ Scoring audit trail: 100% logged

■ ■ Independence audit: Passed 2026-06

PRINCIPLE 8: No Fund Control

■ ■ Payment processors integrated: 0

■ ■ Funds held by Daanaa: \$0

■ ■ User payment complaints: 0

■ ■ Money flow audit: ✓ Clean

PRINCIPLE 9: Explainability

■ ■ Material decisions documented: 100%

■ ■ Methodology history public: Yes

■ ■ DECISIONS.md updated: Monthly

■ ■ Auditor traceability: Full

PRINCIPLE 10: Human Accountability

■ ■ AI-generated scores: 0%

■ ■ AI outputs human-reviewed: 100%

■ ■ Local inference used: 95%

■ ■ Correction rate on AI: 0.8%

PRINCIPLE 11: Strengthening

■ ■ Principle changes this year: 0

■ ■ Silent erosion attempts: 0

■ ■ Re-sign-off compliance: 100%

■ ■ Integrity audit: Quarterly ✓

SUMMARY: 11/11 principles in full compliance

Last audit: 2026-06-08 | Next audit: 2026-07-01

OPERATIONALIZING STEWARDSHIP

In Agent Decision-Making

****Every agent has a Stewardship Profile:****

Agent: Website Discovery Team

Principles_Enforced: [1, 3, 4, 5, 6]

Cannot_Do:

- Prioritize big orgs over small orgs
- Hide orgs from results without explanation
- Make availability decisions based on partnerships
- Surface unverified data as fact

Monthly_Stewardship_Check:

- Are we advantaging large orgs? [No]
- Are small orgs being surfaced? [Yes, 30% discovery rate]
- Any partner pressure to surface specific orgs? [No]

In Code

****Every scoring decision:****

```
def compute_tier(org):
```

```
# IRS data only (Principle 1: Mission)
```

```
score = calculate_from_irs_data(org)
```

```
# Deterministic (Principle 7: Independence)
```

```
assert score == calculate_from_irs_data(org) # Must be same
```

```
# Evidence-based (Principle 3)
```

```
tier = score_to_tier(score)
```

```
tier.metadata = {
```

```
'data_source': 'IRS_990_SOI_2024',
```

```
'methodology_version': 'v2.0',
```

```
'confidence': 0.94, # From peer group size
'peer_group_size': 45000,
'audit_trail': 'logged'
}

# Fair to small orgs (Principle 4)
assert org.peer_group == f'{ntee}_{revenue_band}'

return tier
```

In Operations

****Every morning briefing includes:****

- Stewardship compliance check (any principle breaches?)
- Audit trail summary (decisions made, all logged?)
- Public dashboard update (metrics fresh?)

ESCALATION: When Principles Are Threatened

****If any agent encounters pressure to compromise principles:****

1. PAUSE (do not proceed)
2. LOG (document the pressure, source, timing)
3. ESCALATE (inform CEO Agent immediately)
4. ESCALATE (inform human stakeholder)
5. RESOLVE (CEO Agent + human decide, publicly)
6. DOCUMENT (in Revision Log if principle changes)

****Example escalation:****

- Partnership team: "Partner wants higher ranking"
- → Detected as violation of Principle 7 (Independence)
- → Escalated to CEO Agent
- → CEO Agent confirms: Partnership rejected
- → Logged: "Attempted influence blocked, partner awareness given"

SUCCESS METRICS (Stewardship-Focused)

By end of 2026, Daanaa must prove:

- ■ **Principle 1:** 0 paid placements, 0% mission compromise
- ■ **Principle 2:** 0 server-side donor data, 0 social coercion
- ■ **Principle 3:** 100% of trust signals evidence-based + explained
- ■ **Principle 4:** Equal tier distribution across org sizes, 1K+ hidden gems
- ■ **Principle 5:** 0 shame language, 94%+ org satisfaction
- ■ **Principle 6:** <24h average correction time, 100% disclosed
- ■ **Principle 7:** 0 partner influence, 100% scoring fidelity
- ■ **Principle 8:** \$0 funds held, 0 payment processors
- ■ **Principle 9:** 100% decision documentation, full auditability
- ■ **Principle 10:** 0% AI-generated scores, 100% human oversight
- ■ **Principle 11:** All principle changes public, 0 silent erosion

CONCLUSION

Daanaa's stewardship principles are not suggestions—they are structural constraints embedded in every agent, every decision, every line of code. The organization is designed so that compromising a principle requires deliberate action against the system's own architecture, not just a policy violation.

This is how we earn trust. Not through marketing language, but through systems that make it impossible to do otherwise.

****Stewardship Integration: LIVE****

****Last updated:**** 2026-06-08 13:30 UTC

****Next audit:**** 2026-07-01 (First Friday, Monthly)